

Towards High-Quality LLM-Based Data for French Spontaneous Speech Simplification: an Exo-Refinement Approach

Lucía Ormaechea^{1,2}, Nikos Tsourakis¹, Pierrette Bouillon¹, Benjamin Lecouteux², Didier Schwab²

¹Department of Translation Technology, University of Geneva, Switzerland

²Univ. Grenoble Alpes, CNRS, Grenoble INP, LIG, France

{firstName.lastName}@unige.ch, {firstName.lastName}@univ-grenoble-alpes.fr

Abstract

This study explores the synthetic data generation capabilities of LLMs for French spontaneous speech simplification (S3), a low-resource NLP task. We introduce the exo-refinement approach, which builds on the self-reflect workflow but differs by using separate models for generation and evaluation. To address the limitations of single-model refinement, it integrates external feedback from distinct LLMs as judges, refining outputs based on 3 task-specific dimensions. Comparing expert-based simplifications from gold transcriptions to LLM synthetic outputs, results show that mistral-large outperforms all benchmarked models, including the MUSS baseline, and mistral-small achieves competitive performance with few refinements. SARI results confirm that iterations improve simplicity gain without compromising semantic meaning, as shown by COMET scores. These findings support exo-refinement as a scalable method for synthetic data generation and future S3 model development.

Index Terms: speech simplification, synthetic data generation, large language models

1. Introduction

The generative capabilities of the latest LLMs present a promising avenue for NLP tasks where parallel data is scarce. This has driven research on LLMs as text generation artifacts and have spurred papers exploring prompt engineering techniques [1, 2] and workflows for generating and evaluating synthetic data [3, 4, 5, 6]. These efforts aim to develop methods to produce synthetic data that closely mirrors real-world data in terms of representativity, trustworthiness and diversity.

This paper benchmarks various state-of-the-art LLMs through a data creation pipeline exploring diverse prompting strategies. These are integrated into an iterative generation-evaluation workflow to identify the most effective data creation setup for a low-resource NLP task¹: spontaneous speech simplification (S3) for French. S3 aims at reducing the linguistic complexity of an utterance while preserving its original meaning and grammaticality [7], thus making spontaneous transcripts more accessible for several target audiences (*e.g.*, people with cognitive impairments, non-native speakers, etc.), but also providing an intermediate representation in NLP pipelines involving speech, such as subtitle [8] or pictograph translation [9].

To address the challenges of generating high-quality synthetic simplifications, we build upon the self-refinement generation-evaluation workflow and propose an approach where the tasks of generation and evaluation are handled by separate models, which we coin as *exo-refinement*. This approach enables the infusion of feedback from a distinct LLM evaluator

into the LLM generator, addressing the limitations of relying on a single model for both roles [10]. Our goal is to examine how successive iterations within an exo-refinement approach influence the quality of the generated candidate simplifications, either aligning them more closely to human-crafted versions or rather causing a decline in quality. Identifying the optimal configuration (in terms of model pair, generation strategy and number of iterations) may allow us to scale our data generation workflow and produce a high-quality large-scale synthetic set.

Our contributions are as follows: *i*) we introduce a variation of the self-refinement workflow, termed *exo-refinement*, to generate synthetic simplified outputs from gold transcripts; *ii*) we release a small-scale set of paired spontaneous transcripts with 4 expert-based simplifications; and *iii*) we present the first dedicated solution to automate S3, which far outperforms state-of-the-art text simplification models when applied to speech.

2. Related work

2.1. Spontaneous speech and simplification

S3 is an emerging task and draws inspiration from a well-established NLP field, *i.e.*, text simplification (TS). TS has been widely studied across various written domains such as news articles, Wiki-based content and healthcare documents [11, 12]. Much like its written-centric counterpart, S3 seeks to reduce linguistic complexity while retaining the original meaning and grammaticality. However, most studies have focused on English, leaving simplification under-explored in other languages like French, where sentence-level simplification remains sparsely studied [13, 14].



Figure 1: An utterance extracted from the CEFC corpus [15], comprising its lemmas, POS tags and syntactic information.

While S3 may seem like a mere change in input modality with respect to TS, complexity manifests differently in speech, which varies from written language in ways that extend the transmission mode, impacting lexical, syntactic and morphological aspects [16, 17]. Unlike written texts, which are planned and finalized, spoken texts are an ongoing (and thus unfinished) process. Spontaneous utterances diverge from the sentence structure typical of a written modality, since the discourse thread is less reliant on complete sentences due to cognitive constraints in real-time production. Instead, speech unfolds in

¹Code and prompts are provided in the following repository: <https://github.com/lormaechea/interspeech-2025>.

smaller units (or syntactic windows), which are appended one after the other as it progresses. This leads to disfluent traces like revisions, false starts or reformulations, complicating transcript comprehension and requiring careful treatment during simplification. Real-time construction also impacts information packaging [18], modifying canonical theme-rheme order and emphasis procedures, as seen in French with cleft constructions or dislocated subjects [19]. On-the-fly processing elongates utterances [20] due to performance phenomena that trigger concatenations of elements with a paradigmatic relationship along the syntagmatic axis [21]. This is evident in Figure 1, where paradigmatic supplements weigh down the sentence and hinder its comprehension. Also, context is crucial in simplifying spontaneous speech, given its dialogic nature. Thus, meaning is often not confined to the segment level but derives from the shared interaction between speakers, making the simplification process eventually dependent on prior context.

2.2. LLM-based synthetic data generation

Recent research highlights LLMs as an efficient and practical tool for generating synthetic data to train AI models, especially when human-crafted data is scarce. This is due to two features: *i*) extensive pre-training [22], which equips them with factual and linguistic knowledge to produce coherent human-like text; and *ii*) instruction-following capabilities, which confer models a wide adaptability over textual generation, making them suitable for domain-specific data generation [23, 24], even with limited or no task-specific seed exemplars [25].

LLMs provide extensive customization by allowing users to create and specify attributes in prompts through natural language instructions. However, LLMs lack controllability over outputs, due to their inherent non-determinism and the biases learned during pre-training. This can lead to skewed synthetic datasets [10, 26], factual inaccuracies (*i.e.*, hallucinations) which can ultimately lessen their reliability in generating trustworthy data. Diversity also poses challenges, given that a low one may result in highly monotonous outputs that fail to capture the varied nature of real-world data [22], while a high one can lower accuracy and factual correctness, compromising the trustworthiness of the outputs [5]. Balancing these factors is crucial, as both extremes risk degrading downstream model performance and potentially lead to model collapse².

This underscores the need to distinguish high-quality examples from repetitive ones. While human assessment remains the gold standard, it is notoriously costly in terms of time and budget [27]. This has opened the door to a new paradigm: LLM-assisted evaluation, where machines assess outputs, providing a scalable approach to approximate human preferences [28]. Recent literature presents contrasting views. Optimists see LLM-assisted evaluation as a fast and inexpensive surrogate for human assessment [29], highlighting LLMs' ability to refine outputs through self-refinement, with models being able to rectify errors and produce improved outputs based on their own feedback [3, 4]. Skeptics, however, warn about LLMs' assessment capacities, given their tendency to exhibit biases, often related to specific positions or verbosity [28]. Suboptimality can also be caused by self-bias, leading to overfitting [30], failure to correct mistakes [31] and data diversity harm, especially in smaller models [10]. This suggests the need to integrate external feedback into these workflows to mitigate the impact of such effects.

²See: <https://arxiv.org/abs/2305.17493>.

3. Data

3.1. CEFC speech dataset

S3 remains an underexplored task, compromising the availability of parallel data. In order to evaluate the quality of LLM-based simplifications, we resorted to CEFC (*Corpus d'Étude du Français Contemporain*), a reference corpus for contemporary French spontaneous speech [15], and gathered representative spontaneous utterances. CEFC contains ≈ 270 hours of acoustic data from $\approx 2,500$ speakers across diverse contexts (*e.g.*, lectures, medical consultations, everyday life interactions, among others). Dialogues are distributed across 12 subcorpora and are segmented into audio clips, each representing a conversation.

3.2. Expert-based simplifications from spontaneous speech

The CEFC corpus includes sentence-aligned segments with audio and morphosyntactic annotations, but lacks ground truth simplifications. To address this, we launched a simplification campaign involving 4 expert translators, all native French speakers with backgrounds in translation studies and current positions in the same field. To manage the large source corpus and the high cost of obtaining large-scale manual simplifications, we extracted a small representative subset of CEFC using proportional allocation sampling. The population, namely CEFC, was divided into 12 distinct strata, each corresponding to a subcorpus. From each stratum, we extracted 10 consecutive segments from 2 conversations, resulting in a set of 240 gold transcripts (12 strata $\times$ 2 conversations $\times$ 10 segments) to be simplified by the enlisted professionals (see Table 1).

Table 1: Data overview of the obtained expert simplifications.

	# references	# simplifications
# sentences	240	960
# tokens	3,163	10,083
# types	810	1,107

The simplification task was hosted on the LimeSurvey platform. Prior to participation, all individuals received a consent form outlining the purpose and handling of the data collected. They were asked to provide simpler equivalents for each of the 240 input sentences. In the prompt, which contained detailed instructions, we emphasized that the simplifications should be suitable for a non-native-speaking audience. The sampled sentences were presented in sets of 10, each corresponding to a distinct source conversation. Given the interactive aspect of spontaneous speech, we considered the presence of prior context to be beneficial for eventually disambiguating the next segments.

4. Method

4.1. Model specialization through exo-refinement

Our approach, coined *exo-refinement*, builds on the self-refinement paradigm with a key modification: the *generation* and *evaluation* tasks are conducted by distinct models. This variation allows us to explore the potential of incorporating external feedback in the form of a separate model with advanced reasoning capabilities, addressing issues in prior research that arise from the use of a single LLM for all roles [10, 28]. We investigate whether an iterative generation-evaluation loop, integrating an LLM-assisted evaluation, approximates human intuition and leads to progressively closer alignments with human-crafted simplifications or, conversely, degrades prior outputs.

We thus propose a model specialization of the two above-mentioned tasks:

1. *Generation*, which involves both the initial generation and the further refinements, based on the feedback provided by the evaluation (or judge) LLMs. Generation LLMs were instructed to create outputs using various prompting strategies.
2. *Evaluation*, which involves 3 independent judges that assess 3 core dimensions of simplified texts, with a top evaluator reviewing their scores and making a final decision to accept or reject the proposed simplification.

Table 2: LLMs used for generation and evaluation, including their model size (number of parameters) when available.

LLM family	Role	Model	Size
Mistral	Generators	open-mistral-7b	7B
		open-mistral-nemo-2407	12B
		mistral-small-2409	22B
		mistral-large-2407	123B
	Judge	open-mixtral-8x7b	47B
OpenAI	Generators	gpt-3.5-turbo-0125	<unk>
		gpt-4o-2024-08-06	<unk>
	Judge	gpt-4o-mini-2024-07-18	<unk>

To achieve this, we utilize two LLM families (namely Mistral, mostly open-source, and OpenAI, proprietary). As shown in Table 2, we employ various instruction-following models for generation while keeping the LLM-as-a-judge static. In both LLM clusters, evaluation is performed by an LLM with advanced reasoning capabilities (*open-mixtral-8x7b* and *gpt-4o-mini*). However, they are ranked in a relatively lower level with respect to more recently released LLMs within the same cluster (e.g., *gpt-4o* and *mistral-large*), which are used as generators. All experiments conducted used the respective APIs of Mistral and OpenAI.

4.2. LLMs as generators

Instruction-tuning enables LLMs to generalize to new tasks using natural language prompts. However, given the novelty of the examined task, we hypothesized that instructions alone might not suffice to convey its nuances. We thus explored several prompting techniques to extend the range of candidate simplifications for each input. To generate French outputs with multilingual LLMs, we used English prompting templates, as recent studies [32, 33] show English prompts paired with non-English examples yield better performance, which likely stems from Anglocentric data representation during pre-training. The prompts designed³ were: *i) zero-shot*, where only the task context and generation instructions are provided; *ii) one-shot* and *iii) few-shot*, where one or several task-specific demonstrations are used to implicitly inject human guidance [22] into the prompt; *iv) multi-step*, which couples multiple steps, each dependent on the previous output, guiding the model through the specificity of the task [34]; *v) chain-of-thought*, where human-created rationales encourage the model to break down reasoning into a sequence of steps that lead to a conclusion; *vi) chain-of-thought with self consistency*, which generates multiple reasoning paths for the same query to select the most consistent answer [35]; and *vii) tree-of-thought*, which explores multiple reasoning pathways simultaneously [2], enabling the model to evaluate different approaches in parallel in order to make a final decision.

³All generation prompts followed the same hyperparameter setup: temperature was set to 0.2, max_new_tokens to 256 and top_p to 0.9. The other hyperparameters remained at their default values.

Algorithm 1 Iterative generation-evaluation approach.

Require: conversation, models $\{G, J\}$, prompts $\{p_{gen}, p_{simp}, p_{mean}, p_{gram}, p_{ref}, p_{top}\}$

```

1: for sentence  $x \in$  conversation; do
2:   for iteration  $i \in [0, 4]$ ; do
3:      $y_i, h_i = G(p_{gen}(x, h_{i-1}))$  ▷ Initial generation
4:      $s_i = J(p_{simp}(x, y_i))$  ▷ Starts evaluation
5:      $m_i = J(p_{mean}(x, y_i))$ 
6:      $g_i = J(p_{gram}(y_i))$ 
7:      $top_i = J(p_{top}(x, y_i, s_i, m_i, g_i))$ 
8:     if  $top_i = 0$ ; then ▷  $y_i$  is accepted
9:       return  $y_i$ ; break
10:    else ▷  $y_i$  is rejected
11:      append  $top_i$  to  $fb_i$ 
12:      append  $s_i$  to  $fb_i$  if  $s_i \neq 3$ 
13:      append  $m_i$  to  $fb_i$  if  $m_i \neq 3$ 
14:      append  $g_i$  to  $fb_i$  if  $g_i \neq 0$ 
15:       $y_i, h_i = G(p_{ref}(x, y_i, fb_i, h_i))$ 
16:    end if
17:  end for
18:  return  $y_i$ 
19: end for
20:  $h_i = list()$  ▷ Reset history
```

4.3. LLMs as judges

As discussed earlier, we defined 3 independent LLM-based judges to assess 3 classic dimensions of simplified texts using a criteria-based evaluation, namely:

- *Simplicity gain* and *meaning preservation*, which were evaluated using a 3-point Likert scale (1 being the worst and 3 being the best), and were conceived as ranking tasks. The evaluation was conducted in a reference-guided manner: comparing the simplified outputs against the original utterances.
- *Grammaticality*, which was treated as a reference-free classification task, and assessed using a binary label: 1 for non-grammatical and 0 for grammatical.

For the evaluation templates, we used a uniform structure: *i)* we first provided the judge LLM with a task description; *ii)* we later defined each score value; and *iii)* we outlined the output template, including a numerical score and its reasoning. The decision to proceed to a new refinement is made by a *top evaluator*, who after reviewing the reasoned scores from the judges, provided a final judgment to accept or reject the simplification.

4.4. Generation-evaluation approach

We integrated the above-mentioned templates into the generation and judge models, forming a generation-evaluation loop (see Alg. 1). For each 10-sentence conversation, 7 prompting strategies (p_{gen} ⁴) were used to produce simplifications (y_i) from each input (x). G denotes the generation model and J the evaluation model, which initializes 3 templates: p_{simp} , p_{mean} and p_{gram} . The refinement template is p_{ref} , and p_{top} represents the top evaluator. As discussed in Section 2.1, spontaneous speech is intrinsically conversational, making it potentially highly context-dependent. To address this, we stored the history (h_i) of prompt-response interactions with G for disambiguation or clarification. The history window, limited to 100 entries (10 utterances $\times$ 5 iterations $\times$ 2 user-machine interactions), was reset after each new conversation. No history was stored for the judges, to ensure an isolated LLM-assisted evaluation. The goal was to maintain an objective judgment, and avoid any eventual distortion on the assessment caused by the J model assigning weight to the previous context. Based on the scores and feedback from the 3 judges, which were then sent

⁴For clarity purposes, we used p_{gen} as a global identifier for: $\{p_{zero}, p_{one}, p_{few}, p_{mstep}, p_{cot}, p_{cotsel}, p_{tot}\}$.




Figure 2: Number of outputs obtained per model, COMET results, judge LLMs’ scores and results obtained per prompting strategy.

to the top evaluator, two scenarios can arise. The first scenario occurs when the top evaluator deems that the output should be rejected. In this case, the corresponding feedback (fb_i), along with that from any of the 3 judges assigning a suboptimal score, is collected and sent for regeneration (see l. 10–15 in Alg. 1). This process will continue until the results align with an acceptance response from the top evaluator, or until the maximum of 5 iterations is reached. The second scenario arises when the top evaluator gives a positive review of the reasoned scores from the 3 judges and accepts the output, at which point the refinement process for that sentence stops (see l. 8-9 in Alg. 1).

4.5. Compared methods

S3 lacks task-specific resources that enable its automation using existing models. Consequently, we made use of MUSS [13] as a baseline, the latest reproducible model for sentence-level French TS. MUSS is an unsupervised TS model trained on paraphrasing data, using controllable parameters to adjust length and lexical attributes. We used it in inference mode to produce candidate simplifications from the 240 input utterances.

5. Results and discussion

In this section, we analyze candidate simplifications generated from the exo-refinement approach and compare them to the collected expert simplifications using SARI, BLEU⁵ and COMET⁶ metrics. To this respect, if a candidate simplification is accepted before the 5-iteration limit, the last output is copied into the next ones. This ensures scores are compared using the same number of outputs across all iterations.

Table 3 shows the best results obtained for the 3 metrics across all model configurations. As observed, *mistral-large* stands out as the best performing model, achieving the highest SARI score in the fourth iteration and the highest BLEU score in the first one, while also accomplishing it with one of the lowest number of generated outputs, second only to *mistral-small* (see the leftmost plot of Figure 2). Additionally, it falls only 0.1 points behind the leading model (*i.e.*, *gpt-4o*) in terms of COMET. In light of these results, we also see that the proposed approach consistently surpasses the established baseline by several points, with the sole exception of *open-mistral-7b*, whose best SARI result remains below MUSS. Furthermore, SARI consistently boosts after the initial generation, highlighting the value of an iterative approach for the task at hand.

In contrast, Table 3 reveals the opposite trend for BLEU: the highest scores are obtained in the first iteration, followed by a slight decline in subsequent ones. This indicates that the outputs generated in later cycles diverge lexically from the initial guess, resulting in reduced surface overlap. This pattern reflects the acknowledged inverse SARI-BLEU relationship: as simplicity gain increases, meaning preservation tends to decrease.

⁵We used the HF implementation for both SARI and BLEU metrics.

⁶Model used: Unbabel/XCOMET-XL.

Table 3: Extracted best SARI, BLEU and COMET results for each model configuration, including iteration and prompt used.

model	SARI (it)	gain	prompt	BLEU (it)	gain	prompt	COMET (it)	prompt
MUSS	50.75	—	—	41.66	—	—	64.58	—
open-mistral-7b	49.46 (2)	+1.26	m-step	45.50 (1)	+3.69	tot	90.78 (5)	m-step
open-mistral-nemo	54.59 (5)	+0.26	0-shot	47.86 (1)	+1.98	0-shot	93.44 (2)	m-step
mistral-small-2409	58.47 (2)	+0.66	0-shot	56.80 (1)	+1.31	0-shot	93.92 (5)	m-step
mistral-large-2407	60.22 (4)	+0.52	1-shot	60.89 (1)	+1.88	0-shot	94.19 (3)	1-shot
gpt-3.5-turbo-0125	58.45 (5)	+1.72	tot	56.14 (5)	+0.62	0-shot	93.74 (3)	tot
gpt-4o-2024-08-06	58.59 (2)	+1.76	0-shot	55.08 (1)	+1.13	0-shot	94.29 (1)	m-step

However, as shown in the second plot in Figure 2, COMET remains consistently stable across all iterations, indicating that while lexical surface changes, the underlying semantic meaning remains intact. This reinforces the idea that an iterative exo-refinement approach contributes to the generation of more suited outputs for S3. We also analyzed the performance of the judge LLMs for simplicity gain and meaning retention on sentences that reached the maximum of iterations (see yellow slices in the leftmost plot of Figure 2), which correspond to the most challenging cases for the proposed framework. The heatmap in Figure 2 illustrates the average score for each dimension across iterations. Consistently, *mistral-small* and *mistral-large* showed superior performance within their model family: a positive indication of their effectiveness on challenging inputs. However, further investigation is required to contrast these results with OpenAI models, which use a different LLM-as-a-judge, and criteria applied may be stricter or more lenient as applicable. Regarding the last plot in Figure 2, it is interesting to note that, despite our intention to explore a diverse range of prompting strategies, the default *zero-shot* template mostly delivered the best performance in terms of SARI and BLEU.

6. Conclusions

In this paper, we present an iterative generation-evaluation approach to examine the simplification capabilities of a set of OpenAI and Mistral LLMs. Outputs were compared to expert simplifications, and results showed that *mistral-large* outperformed all other models tested, including the MUSS baseline. Moreover, *mistral-small* required the fewest refinements and reached very competitive performance compared to the other LLMs. Techniques such as *zero-shot* and *multi-step* yielded the highest results in terms of SARI, BLEU and COMET, demonstrating the potential to scale this approach for large-scale data generation and the further development of an S3 downstream model.

We acknowledge several limitations in this study. We benchmarked only a subset of recent LLMs, and experiments used gold transcriptions rather than speech signals. As further work, we plan to integrate this approach with an ASR system. Also, experiments were conducted on a small corpus, given the limited S3 resources and the need to balance workload with quality in the expert simplification task. Finally, while automatic scores indirectly correlate with human correctness via ground-truth simplifications, LLM-as-a-judge scores may require human calibration for a more comprehensive assessment.

7. Acknowledgements

This work is part of the PROPICTO project, funded by the Swiss National Science Foundation (N°197864) and the French National Research Agency (ANR-20-CE93-0005).

8. References

- [1] T. B. Brown *et al.*, “Language models are few-shot learners,” in *Advances in Neural Information Processing Systems*, H. Larochelle *et al.*, Eds., vol. 33, 2020, pp. 1877–1901. [Online]. Available: <https://api.semanticscholar.org/CorpusID:218971783>
- [2] S. Yao *et al.*, “Tree of thoughts: Deliberate problem solving with large language models,” in *International Conference on Neural Information Processing Systems*, 2023. [Online]. Available: <https://dl.acm.org/doi/abs/10.5555/3666122.3666639>
- [3] A. Madaan *et al.*, “SELF-REFINE: iterative refinement with self-feedback,” in *Advances in Neural Information Processing Systems*, 2023. [Online]. Available: <https://api.semanticscholar.org/CorpusID:257900871>
- [4] N. Shinn *et al.*, “Reflexion: Language agents with verbal reinforcement learning,” in *International Conference on Neural Information Processing Systems*, 2023. [Online]. Available: <https://dl.acm.org/doi/10.5555/3666122.3666499>
- [5] J. Chung *et al.*, “Increasing diversity while maintaining accuracy: Text data generation with large language models and human interventions,” in *Annual Meeting of ACL*, A. Rogers *et al.*, Eds., 2023, pp. 575–593. [Online]. Available: <https://aclanthology.org/2023.acl-long.34>
- [6] Y. Liu *et al.*, “Calibrating LLM-based evaluator,” in *LREC-COLING*, 2024, pp. 2638–2656. [Online]. Available: <https://aclanthology.org/2024.lrec-main.237>
- [7] C. Horn *et al.*, “Learning a lexical simplifier using Wikipedia,” in *Annual Meeting of ACL*, K. Toutanova and H. Wu, Eds., 2014, pp. 458–463. [Online]. Available: <https://aclanthology.org/P14-2075>
- [8] S. Mehta *et al.*, “Simplify-then-translate: Automatic pre-processing for black-box translation,” in *AAAI Conference on Artificial Intelligence*, 2020. [Online]. Available: <https://ojs.aaai.org/index.php/AAAI/article/view/6369>
- [9] C. Macaire *et al.*, “A multimodal French corpus of aligned speech, text, and pictogram sequences for speech-to-pictogram machine translation,” in *LREC-COLING*, 2024, pp. 839–849. [Online]. Available: <https://aclanthology.org/2024.lrec-main.76>
- [10] W. Xu *et al.*, “Pride and prejudice: LLM amplifies self-bias in self-refinement,” in *Annual Meeting of ACL*, 2024, pp. 15 474–15 492. [Online]. Available: <https://aclanthology.org/2024.acl-long.826>
- [11] —, “Problems in current text simplification research: New data can help,” *Transactions of ACL*, vol. 3, pp. 283–297, 2015. [Online]. Available: <https://direct.mit.edu/tacl/article/43283>
- [12] X. Zhang and M. Lapata, “Sentence simplification with deep reinforcement learning,” in *EMNLP*, 2017, pp. 584–594. [Online]. Available: <https://aclanthology.org/D17-1062>
- [13] L. Martin *et al.*, “MUSS: Multilingual unsupervised sentence simplification by mining paraphrases,” in *LREC*, 2022, pp. 1651–1664. [Online]. Available: <https://aclanthology.org/2022.lrec-1.176>
- [14] L. Ormaechea and N. Tsourakis, “Automatic text simplification for French: Model fine-tuning for simplicity assessment and simpler text generation,” *IJST*, vol. 27, p. 957–976, 2024. [Online]. Available: <https://link.springer.com/article/10.1007/s10772-024-10146-0>
- [15] C. Benzitoun *et al.*, “Le projet ORFÉO : un corpus d’études pour le français contemporain,” *Corpus n° 15*, pp. 91–114, 2016. [Online]. Available: <https://journals.openedition.org/corpus/2936>
- [16] A. Caines *et al.*, “Parsing transcripts of speech,” in *Workshop on Speech-Centric Natural Language Processing*, N. Ruiz and S. Bangalore, Eds., 2017, pp. 27–36. [Online]. Available: <https://aclanthology.org/W17-4604>
- [17] L. Ormaechea *et al.*, “Simplification strategies in French spontaneous speech,” in *Proceedings of the DeTermIt Workshop @ LREC-COLING*, 2024, pp. 90–102. [Online]. Available: <https://aclanthology.org/2024.determin-1.9/>
- [18] M. A. K. Halliday, *Spoken and written language*. Oxford University Press, 1985.
- [19] C. Blanche-Benveniste, *Approches de la langue parlée en français*. Ophrys, 1997.
- [20] T. Bazillon *et al.*, “La parole spontanée : transcription et traitement,” in *TAL*, B. Laks and N. Nguyen, Eds., vol. 49, 2008, pp. 47–76. [Online]. Available: <https://aclanthology.org/2008.tal-3.3>
- [21] D. Luzzati, “Rhétorique et description de l’oral,” *VERBA*, vol. 25, pp. 7–30, 1998.
- [22] L. Long *et al.*, “On LLMs-driven synthetic data generation, curation, and evaluation: A survey,” in *Findings of ACL*, L.-W. Ku *et al.*, Eds., 2024, pp. 11 065–11 082. [Online]. Available: <https://aclanthology.org/2024.findings-acl.658>
- [23] O. Honovich *et al.*, “Unnatural instructions: Tuning language models with (almost) no human labor,” in *Annual Meeting of ACL*, A. Rogers *et al.*, Eds., 2023, pp. 14 409–14 428. [Online]. Available: <https://aclanthology.org/2023.acl-long.806>
- [24] Y. Wang *et al.*, “Self-Instruct: Aligning language models with self-generated instructions,” in *Annual Meeting of ACL*, A. Rogers *et al.*, Eds., 2023, pp. 13 484–13 508. [Online]. Available: <https://aclanthology.org/2023.acl-long.754>
- [25] A. Rebmann *et al.*, “Evaluating the ability of LLMs to solve semantics-aware process mining tasks,” *International Conference on Process Mining*, pp. 9–16, 2024.
- [26] D. Shanley *et al.*, “Getting real about synthetic data ethics,” *EMBO reports*, vol. 25, no. 5, pp. 2152–2155, 2024.
- [27] S. Shankar *et al.*, “Who validates the validators? aligning LLM-assisted evaluation of LLM outputs with human preferences,” in *Annual ACM Symposium on User Interface Software and Technology*, 2024. [Online]. Available: <https://doi.org/10.1145/3654777.3676450>
- [28] L. Zheng *et al.*, “Judging LLM-as-a-judge with MT-bench and Chatbot Arena,” in *Advances in Neural Information Processing Systems*, vol. 36, 2023, pp. 46 595–46 623. [Online]. Available: <https://dl.acm.org/doi/10.5555/3666122.3668142>
- [29] F. Gilardi *et al.*, “ChatGPT outperforms crowd workers for text-annotation tasks,” *Proceedings of the National Academy of Sciences*, vol. 120, no. 30, 2023. [Online]. Available: <http://dx.doi.org/10.1073/pnas.2305016120>
- [30] Y. Liu *et al.*, “G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment,” in *EMNLP*, H. Bouamor *et al.*, Eds., 2023, pp. 2511–2522. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.153>
- [31] J. Huang *et al.*, “Large language models cannot self-correct reasoning yet,” in *International Conference on Learning Representations*, 2024. [Online]. Available: <https://iclr.cc/virtual/2024/poster/18956>
- [32] F. Shi *et al.*, “Language models are multilingual chain-of-thought reasoners,” in *International Conference on Learning Representations*, 2023. [Online]. Available: <https://iclr.cc/virtual/2023/poster/10960>
- [33] B. Zhang *et al.*, “Prompting large language model for machine translation: A case study,” in *International Conference on Machine Learning*, 2023, pp. 41 092–41 110. [Online]. Available: <https://dl.acm.org/doi/abs/10.5555/3618408.3620130>
- [34] P. Liu *et al.*, “Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing,” *ACM Comput. Surv.*, vol. 55, no. 9, 2023. [Online]. Available: <https://doi.org/10.1145/3560815>
- [35] X. Wang *et al.*, “Self-consistency improves chain of thought reasoning in language models,” in *International Conference on Learning Representations*, 2023. [Online]. Available: <https://iclr.cc/virtual/2023/poster/11718>